

# Lab9: Halloween Mini Project

Duy An Le (PID: A16400411)

## Import Data

```
candy_df <- read.csv("candy-data.txt", row.names=1)
head(candy_df)
```

	chocolate	fruity	caramel	peanutyalmondy	nougat	crispedricewafer
100 Grand	1	0	1	0	0	1
3 Musketeers	1	0	0	0	1	0
One dime	0	0	0	0	0	0
One quarter	0	0	0	0	0	0
Air Heads	0	1	0	0	0	0
Almond Joy	1	0	0	1	0	0

	hard	bar	pluribus	sugarpercent	pricepercent	winpercent
100 Grand	0	1	0	0.732	0.860	66.97173
3 Musketeers	0	1	0	0.604	0.511	67.60294
One dime	0	0	0	0.011	0.116	32.26109
One quarter	0	0	0	0.011	0.511	46.11650
Air Heads	0	0	0	0.906	0.511	52.34146
Almond Joy	0	1	0	0.465	0.767	50.34755

## Data Exploration

Q.1 How many different candies are in the dataset?

There are 85 candies in this dataset.

Q.2 How many fruity candies are in the dataset?

There are 38 fruity candies in the dataset.

My favorite candy is Kit Kat.

```
candy_df["Kit Kat",]
```

```

      chocolate fruity caramel peanutyalmondy nougat crispedricewafer hard
Kit Kat      1      0      0              0      0              1      0
      bar pluribus sugarpercent pricepercent winpercent
Kit Kat      1      0      0.313      0.511      76.7686

```

Q.3/Q.4 What is your favorite candy and what is its winpercent? What is Kit Kat's winpercent?

Kit Kat has a winpercent of 76.7686

Q.5 What is "Tootsie Roll Snack Bars" winpercent?

Tootsie Roll Snack Bars have a winpercent of 49.653503

```
skimr::skim(candy_df)
```

Table 1: Data summary

Name	candy_df
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12
Group variables	None

### Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmondy	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q.6 Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?


The sugar percent, price percent, and winpercent seem to be on a different scale

Q.7 What do you think a zero and one represent for the candy\$chocolate column?

Zeros represent the candy is not chocolate based and 1 means it is.

Q.8 Plot a histogram of winpercent values

```
library(ggplot2)
ggplot(candy_df) + aes(x=winpercent) + geom_histogram(bins=7)
```



Q.9 Is the distribution of winpercent values symmetrical?

The distribution is not symmetrical

Q.10 Is the center above or below 50%?

Below 50%

Q.11 On average is chocolate candy higher or lower ranked than fruit candy?

```
chocolate_win <- mean(candy_df[candy_df$chocolate == 1,]$winpercent)
fruit_win <- mean(candy_df[candy_df$fruity == 1,]$winpercent)
```

Chocolate is ranked higher with an average winpercent 60.9215294 while fruity candies are only ranked with an average winpercent 44.1197414

Q.12 Is this difference statistically significant

```
chocolate_df <- candy_df[candy_df$chocolate == 1,]
fruity_df <- candy_df[candy_df$fruity == 1,]

ttest <- t.test(chocolate_df$winpercent, fruity_df$winpercent)
ttest
```

Welch Two Sample t-test

```
data: chocolate_df$winpercent and fruity_df$winpercent
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
 60.92153  44.11974
```

Since the t-test had a p-value  $2.8713778 \times 10^{-8}$  the difference is statistically significant.

Q.13 What are the five least liked candy types in this set?

```
candy_df <- candy_df[order(candy_df$winpercent), ]
row.names(head(candy_df, n=5))
```

```
[1] "Nik L Nip"           "Boston Baked Beans" "Chiclets"
[4] "Super Bubble"       "Jawbusters"
```


Q.14 What are the top 5 all time favorite candy types out of this set?

```
row.names(tail(candy_df, n=5))
```

```
[1] "Snickers"           "Kit Kat"
[3] "Twix"              "Reese's Miniatures"
[5] "Reese's Peanut Butter cup"
```

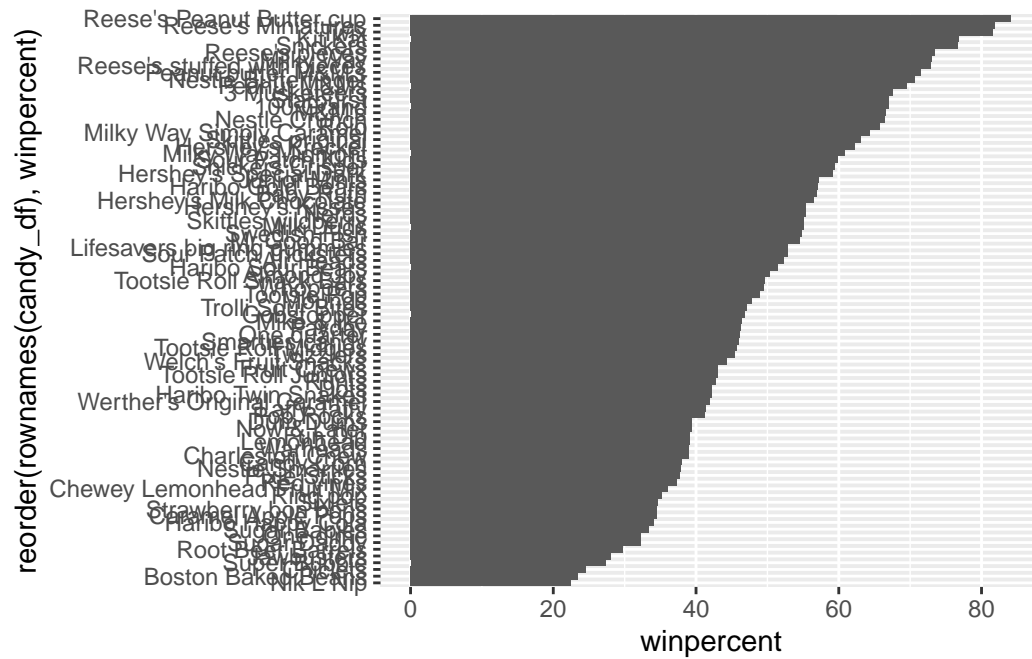
Q.15 Make a first barplot of candy ranking based on winpercent values.

```
ggplot(candy_df) + aes(winpercent, rownames(candy_df)) + geom_col()
```




Q.16 This is quite ugly, use the reorder() function to get the bars sorted by winpercent?

```
ggplot(candy_df) + aes(winpercent, reorder(rownames(candy_df), winpercent)) + geom_col()
```



Adjust figure height:

Saving 5.5 x 10 in image




Add some color:

```
plot.cols <- rep("green3", nrow(candy_df))
plot.cols[as.logical(candy_df$chocolate)] = "chocolate4"
plot.cols[as.logical(candy_df$fruity)] = "maroon"
plot.cols[as.logical(candy_df$caramel)] = "gold2"

ggplot(candy_df) + aes(winpercent, reorder(rownames(candy_df), winpercent)) + geom_col(fill=plot.cols)
```





Q.17 What is the worst ranked chocolate candy?

Sixlets are the worst ranked chocolate candy.

Q.18 What is the best ranked fruity candy?


The best ranked fruity candy is starbursts

Plot winpercent vs. pricepercent

```
library(ggrepel)

ggplot(candy_df) +
  aes(winpercent, pricepercent, label=rownames(candy_df)) +
  geom_point(col=plot.cols) +
  geom_text_repel(col=plot.cols, size=3.3, max.overlaps = 5)
```

Warning: ggrepel: 65 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Q.19 Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

Reese's Miniatures offer the most bang for your buck.

Q.20 What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
temp <- candy_df[order(candy_df$pricepercent, decreasing=T), ]
head(temp, n=5)
```

	chocolate	fruity	caramel	peanutyalmondy	nougat
Nik L Nip	0	1	0	0	0
Nestle Smarties	1	0	0	0	0
Ring pop	0	1	0	0	0
Mr Good Bar	1	0	0	1	0
Hershey's Milk Chocolate	1	0	0	0	0

	crispedricewafer	hard	bar	pluribus	sugarpercent
Nik L Nip	0	0	0	1	0.197
Nestle Smarties	0	0	0	1	0.267
Ring pop	0	1	0	0	0.732
Mr Good Bar	0	0	1	0	0.313
Hershey's Milk Chocolate	0	0	1	0	0.430

	pricepercent	winpercent
Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719
Ring pop	0.965	35.29076
Mr Good Bar	0.918	54.52645
Hershey's Milk Chocolate	0.918	56.49050

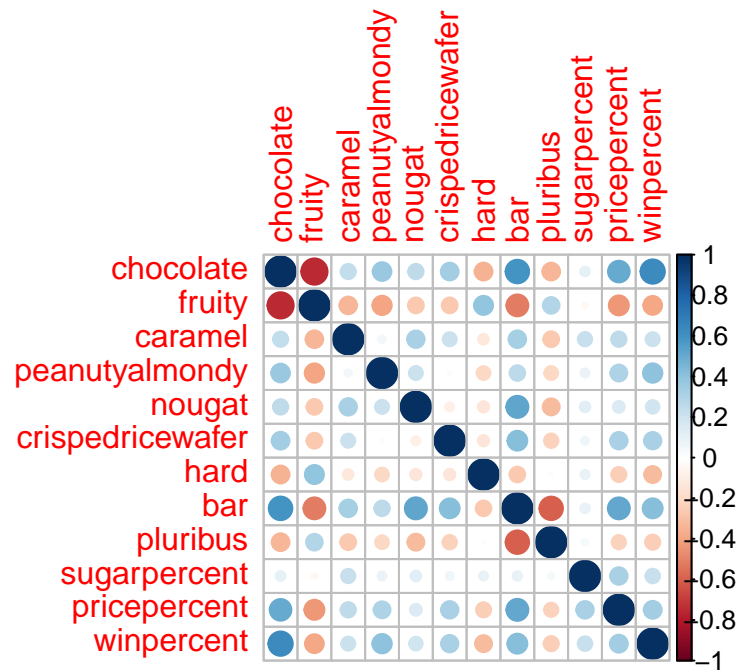
Nik L Nip are the least popular of the most expensive candies.

## Correlation

```
library(corrplot)
```

corrplot 0.92 loaded

```
cij <- cor(candy_df)
corrplot(cij)
```



Q.22 Which two variables are anti-correlated?

Chocolate and fruit are anti-correlated

Q.23 Which two variables are most strongly correlated?

Chocolate and win percent are positively correlated.

## PCA

```
pca <- prcomp(candy_df, scale.=T)
summary(pca)
```

Importance of components:


	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369

	PC8	PC9	PC10	PC11	PC12
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000

Plot PC1 vs PC2


```
plot(pca$x[,1:2], col=plot.cols, pch=16)
```



Use ggplot


```
ggplot(pca$x) + aes(x=pca$x[,1], y=pca$x[,2], size=candy_df$winpercent/100) + geom_point(candy_df) +  
  labs(title = "Halloween Candy PCA space") + xlab("PC1") + ylab("PC2") +  
  geom_text_repel(col=plot.cols, size=3.3, max.overlaps=7, label=row.names(candy_df))
```

Warning: ggrepel: 68 unlabeled data points (too many overlaps). Consider increasing max.overlaps



How do the original variables contribute to PCs? For this we look at the loadings component of our results object ie. the `pca$rotation` object.


```
par(mar=c(8,4,2,2))
barplot(pca$rotation[,1], las=2, ylab="PC1 Contribution")
```



Or with ggplot

```
res <- pca$rotation

ggplot(res) + aes(PC1, rownames(res)) + geom_col() + ylab("PC1 contribution")
```



Q.24 What original variables are picked up strongly by PC1 in the negative direction? Do these make sense to you?

Fruit, pluribus, and hard are all picked up in the negative direction and these make sense based on the correlation analysis done previously.